

## Creating/Altering Tables

```
CREATE TABLE <TableName> (  
    <columnName> <columnType>, ...,  
    CONSTRAINT <constraintName> <constraintType> (<columnName(s)>), ....);
```

Column Types: Number(#), varchar2(#), date

Adding via Alter: ALTER TABLE <TableName> ADD <valid column or constraint>;

## Constraint Examples

```
CONSTRAINT Table_PK PRIMARY KEY (keyColumn);  
CONSTRAINT Table_FKCol_FK FOREIGN KEY (FKCol) REFERENCES OtherTable (key);  
CONSTRAINT Table_UQ UNIQUE (uniqueColOne, uniqueColTwo);  
CONSTRAINT Table_SetInVal CHECK (Set IN ('val1', 'val2'));  
CONSTRAINT Table_ValNotNull CHECK (Val IS NOT NULL);  
    Also: ALTER TABLE <TableName> MODIFY (<colName> NOT NULL);  
CONSTRAINT Table_StringLikeVal CHECK (StringCol LIKE 'pattern');  
CONSTRAINT Table_RangeVal CHECK (Range < upperLimit AND Range > lowerLimit);
```

## Nulls

- 'Null' means UNKNOWN → causes undefined behavior in some operations
- Null groups into its own group
- Ignored in all aggregate functions except COUNT(\*)
- NVL(exp1, exp2) = exp1 if exp1 is not null, else exp2

## DML

```
INSERT INTO <TableName> [{<columnName>}] VALUES ({<columnValues>});  
DELETE FROM <TableName> WHERE <boolExp>;
```

```
UPDATE <TableName> SET <columnName> = <valueExpression> [WHERE <boolExp>];
```

Update on a condition:

```
UPDATE Table SET col = 100 WHERE condition = 'yes';
```

Update from another table:

```
UPDATE ATable A SET col = (  
    SELECT B.value FROM BTable B WHERE A.id = B.id)  
WHERE id IN (SELECT B.id FROM BTable B);
```

```
SELECT [DISTINCT] _ FROM _ [WHERE _] [GROUP BY _ [HAVING _]] [ORDER BY _]
```

Execution Order: FROM/WHERE/GROUP BY/HAVING/ORDER BY/SELECT

Aggregation Operators: COUNT(\*), COUNT(x), MIN(x), MAX(x), SUM(x), AVG(x)

Non-Aggregated columns in the SELECT clause must be in the GROUP BY clause

Renaming in the SELECT Clause: <colName or aggregation or expression> AS <newName>

Formatting Numbers: to\_char(<number>, 990.99)

String concatenation: 'string1' || 'string2'

# OPTIMIZATION OF RELATIONAL ALGEBRA

--> Do Selection first, then project, only join when you have to!

- ① JOIN v. CART. PRODUCT  
- joins are better

$$R \bowtie S = \sigma_c(R \times S)$$

- ② PUSH SELECTION DOWN  
⇒ DO SELECTION AS EARLY AS POSSIBLE

- it reduces the size of the data

$$\sigma_c(\pi_L(R)) \leftrightarrow \pi_L(\sigma_c(R))$$

$$\sigma_{R \times S}(R \bowtie_{R.a=S.b} S) \leftrightarrow \sigma_{R \times S}(R) \bowtie_{R.a=S.b} S$$

- ③ AVOID UNNECESSARY JOINS

find customers having account balances below 100 and loans above 10000

$$R1 \leftarrow \pi_{\text{CustName}}(\text{Depositor} \bowtie \pi_{\text{AccountNum}}(\sigma_{\text{balance} < 100}(\text{Account})))$$

$$R2 \leftarrow \pi_{\text{CustName}}(\text{Borrower} \bowtie \pi_{\text{LoanNum}}(\sigma_{\text{Amount} > 10000}(\text{Loan})))$$

$$\text{Result} \leftarrow R1 \cap R2$$

⇒ Much better than trying to join 4 tables!

## RELATIONSHIPS AMONG OPERATORS

- ① JOIN ↔ CARTESIAN PRODUCT + SELECT  
 $R \bowtie S = \sigma_c(R \times S)$
- ② SELECTION is commutative  
 $\sigma_{c_2}(\sigma_{c_1}(R)) = \sigma_{c_1}(\sigma_{c_2}(R)) = \sigma_{c_1 \text{ AND } c_2}(R)$
- ③ ORDER BETWEEN SELECTION & PROJECTION  
 $\sigma_c(\pi_L(R)) \rightarrow \pi_L(\sigma_c(R))$   
 $\pi_L(\sigma_c(R)) \rightarrow \sigma_c(\pi_L(R))$  \* ONLY if L has all cols needs for C
- ④ JOIN is COMMUTATIVE  
 $R \bowtie S = S \bowtie R$
- ⑤ ORDER BETWEEN SELECTION & JOIN  
 $\sigma_{c_1}(R \bowtie_{c_2} S) \rightarrow (\sigma_{c_1}(R) \bowtie_{c_2} S)$

- ① Set operations: UNION, INTERSECTION, DIFFERENCE



- ② PROJECTION  $\pi_L(R)$ : SELECT clause - chooses some columns to use  
L is a list consisting of col name, col renames (eg A as B), or expressions (eg A+B as Z)

- ③ SELECTION  $\sigma_c(R)$ : WHERE clause - choose some tuples to use  
C is a conditional to apply to each tuple in R

- ④ COMBINING TABLES: FROM clause - choose which tables to use & how to join
- ① CROSS-PRODUCT  $A \times B$ : pairs each  $a \in A$  with each  $b \in B$  ⇒ makes huge tables
  - ② NATURAL JOIN  $A \bowtie B$ : pairs each matching attribute in matching columns
  - ③ THETA JOIN  $A \bowtie_c B$ : pairs each  $a \in A$  with each  $b \in B$  if  $c(a,b)$  holds true
  - ④ OUTER JOIN  $A \bowtie_{\text{or}} B$ : pairs according to theta join, then pass the dangling tuples with ⊥

A:	X	Y
	0	1
	1	0

B:	X	C
	1	0
	1	1

$A \times B$ :	X	Y	B	X	C
	0	1	1	0	0
	1	0	1	0	0
	0	1	1	1	1
	1	0	1	1	1

$A \bowtie B$ :	X	Y	C
	1	0	0
	1	0	1

$A \bowtie_{Y=C} B$ :	X	Y	C
	0	1	1
	1	0	1

$A \bowtie_{\text{or}} B$ :	X	Y	C
	1	0	0
	1	0	1
	0	1	⊥

- ④ RENAMING:  $\rho_{s(A_1, A_2, \dots, A_n)}(R)$ : AS operator

- ⑤ DUPLICATE ELIMINATION:  $\delta(R)$ : DISTINCT operator - returns R with one copy of each tuple in R  
⇒ Turns a bag into a set

- ⑥ SORTING:  $\tau_L(R)$ : ORDER BY clause

- L is a list of fields to sort by, where ties are broken by fields later in the list

- ⑦ AGGREGATION & GROUPING:  $\gamma_L(R)$ : GROUP BY clause

- L is a list of GROUPING ATTRIBUTES (attributes ∈ R) and AGGREGATED ATTRIBUTES (operator applied to cols of R)


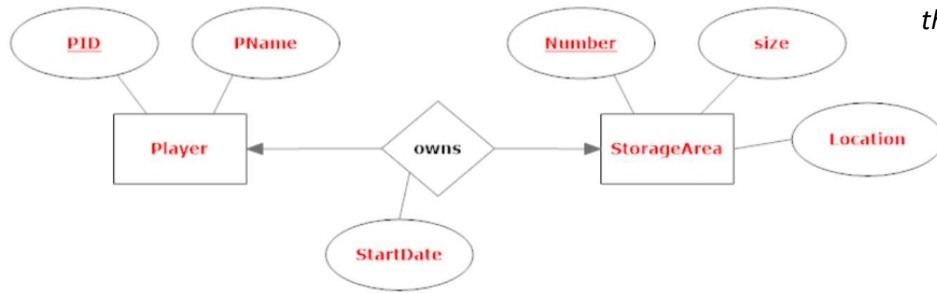
$\gamma_{\text{activity, sum(hours)} \rightarrow \text{timespent}}(\text{HoursLog})$

- executing  $\gamma_L(R)$ :
- ① Partition R into groups, where each group has tuples with a distinct assignment of the grouping attributes  
→ if no grouping attrs specified, R is one group
  - ② For each group, produce one tuple consisting of:
    - the grouping attributes for that group
    - the aggregations over the tuples in that group

# ERDs: Basic Rules


**1:1 or 1:M (with no relationship keys)**  
relationships become part of one of the entity tables

- If 1:1, PK of one side is copied to the other as a FK
- If 1:M, PK of the "one" side is copied to the "many" side as a FK
- Any relationship attributes go on the side with the FK





**1:1 or 1:M (with relationship keys)**  
relationships map to a separate table

- Relationship maps to a table with its PK and Attributes, plus the PK from the "many" side
  - Does NOT get the PK from the "one" side
- "many" side has a FK that references the "one" side's PK
- In a recursive relationship, the entity contains a FK to itself



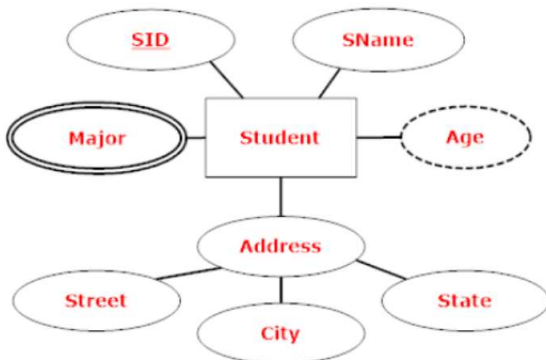

**M:M or Multi-way relationships** map to a separate table

- Relationship table's PK = Keys from each entity + Relationship keys



**Weak Entity Sets** become a separate table

- Contains all its attributes and FKs to the PKs of identifying entity sets
- PK = ID'ing Keys from the identifying Entity sets + Discriminator
- The total, 1:M supporting relationship is not mapped



**Composite Attributes** only use the 2nd level attributes as columns  
**Derived Attributes** map as-is (and enforce via triggers)

**Multi-Valued Attributes** become a separate table

PK = Attribute itself + PK of main entity set

# ERDs: IsA

## A: Relation for each Entity Set

- Redundancy is inherent to the design (a student is in both the Students and People tables)
- Need multiple tables to complete a record for a student or employee

### A1: Partial/Overlapping

Person (SSN, Name, DOB)  
Employee (SSN, Department, Salary)  
FK Employee(SSN) Ref Person(SSN)  
Student (SSN, GPA, StartDate)  
FK Students(SSN) Ref Person(SSN)

### A2: Disjoint/Total

Person (SSN, Role, Name, DOB)  
Unique (SSN, Role)  
Role in {'Student', 'Employee'}  
Employee (SSN, Role, Department, Salary)  
Employee.Role = 'Employee'  
Student (SSN, Role, GPA, StartDate)  
Role = 'Student'

## B: One Big Table

- Full of nulls
- Less joins

### B1: Partial/Overlapping

Person (SSN, Name, DOB, Department, Salary, GPA, StartDate)

### B2: Disjoint/Total

- Need triggers to enforce attribute values based on role

Person (SSN, Name, DOB, Role, Department, Salary, GPA, StartDate)

Unique (SSN, Role)

Role in {'Student', 'Employee'}

## C: Relations only for Specialization

- Cannot be used for Partial, but good for total
- If overlapping, need to update both tables on updates
- If disjoint, need triggers to ensure not in both tables
- Best for total disjoint

Employee (SSN, Name, DOB, Department, Salary)

Student (SSN, Name, DOB, GPA, StartDate)

## D: Relation for Every Combination

- Could get out of hand with lots of overlapping specializations
- Needs triggers to ensure a record is only in one table
- Good for overlapping relationships
- If Partial, need an additional table:  
Person (SSN, Name, DOB)
- Probably simpler to use A2

Employee (SSN, Name, DOB, Department, Salary)

Student (SSN, Name, DOB, GPA, StartDate)

StudentEmployee (SSN, Name, DOB, Department, Salary, GPA, StartDate)

```
create table Person (  
    PID number(3),  
    Role varchar2(10),  
    name varchar2(30),  
    DoB date,  
    constraint Person_pk primary key (PID),  
    constraint Person_un unique (PID, Role),  
    constraint PersonRoleVal check (Role in ('Student', 'Employee'))  
);  
  
create table Student (  
    PID number(3),  
    Role varchar2(30) default 'Student' not null,  
    GPA number(2,1),  
    constraint Student_pk primary key (PID),  
    constraint StudentRoleVal check (Role in ('Student')),  
    constraint Student_fk foreign key (PID, Role) references Person (PID, Role)  
);  
  
create table Employee (  
    PID number(3),  
    Role varchar2(30) default 'Employee' not null,  
    Salary number(6),  
    constraint Employee_pk primary key (PID),  
    constraint EmployeeRoleVal check (Role in ('Employee')),  
    constraint Employee_fk foreign key (PID, Role) references Person (PID, Role)  
);
```

## Views

CREATE [OR REPLACE] VIEW <vName> AS <Query>;

## Triggers

Exec Order:

1. BEFORE (Statement-Level);
2.  $\forall$  affected records: (a) BEFORE (Row-Level); (b) Event (Row; (c) AFTER (Row);
3. AFTER (Statement);

Variables

- Declaration like normal but with a semicolon: <varName> <varType>;
- Can also declare Cursors: CURSOR <cName> IS <query>;
  - Parameterized Cursor: CURSOR <cName> (<param> <type>, ...) IS...;
- Set variables from a table by SELECT <column[s]> INTO <varName[s]>...;
  - System Variables: SELECT sysdate INTO temp FROM Dual;

Code Body

- If statements are explicit: IF (<condition>) THEN <code> END IF;
- Looping through a cursor: FOR row IN cName LOOP <code> END LOOP;
- In Row-Level, may get :new and :old variables (depending on triggering operation)
- Output to console: dbms\_output.put\_line('message');
- Raise Error: RAISE\_APPLICATION\_ERROR(-20001, 'errMessage');

CREATE [OR REPLACE] TRIGGER <TriggerName>

[BEFORE | AFTER] [INSERT | UPDATE [OF <columnName>]] ON <TableName>

[FOR EACH ROW] -- "FOR EACH STATEMENT" implicit if omitted

[DECLARE

    <Declaration>; ...]

BEGIN

    <PL/SQL Code>

END; /

## Procedures & Functions

- Procedures can't output except via output parameters
- Invoke procedures: EXEC <pName>(<params>); (May need to SET serveroutput on;)
- Invoke Functions anywhere with <fName>(<params>), including in WHERE clause
- Access variables declared with the name of procedure or function: <name>.<varName>

CREATE [OR REPLACE] [PROCEDURE | FUNCTION] <name>

    [(<paramName> [IN | OUT] <paramType>, ...)]

    [RETURN <returnType>] -- only if this is a function

    IS

    [<varDeclaration>;...]

BEGIN

    <PL/SQL Code>

END <pName>; /

```

import java.sql.*;
public class OTest {
    // These are for the Database you are connecting to
    private static final String USERID = "USERID"; // also set PASSWORD
    private static final String DB_SERVER = "jdbc:oracle:thin:@oracle.wpi.edu:1521:orcl"
    public static void main(String[] args) {
        try { Class.forName("oracle.jdbc.driver.OracleDriver");
        } catch (ClassNotFoundException e) { // Driver not installed...
            e.printStackTrace(); return; }
        Connection conn = null;
        try { conn = DriverManager.getConnection(DB_SERVER, USERID, PASSWORD);
        } catch (SQLException e) { // Connection Failed...
            return; }
        try {
            Statement stmt = conn.createStatement(); // Basic way (How we did in class)
            String str = "SELECT * FROM TableName";
            ResultSet rset = stmt.executeQuery(str);
            // Process the results
            int custID = 0; String custName = ""; String city = ""; int age = 0;
            while (rset.next()) { // For each row that was returned...
                custID = rset.getInt("id"); // also getString, getDate, ...
            }
            rset.close(); stmt.close(); // Close unneeded resources in this order

            Scanner reader = new Scanner(System.in); // Get User Input
            System.out.println("Enter parameter: ");
            String parameter = reader.nextLine();
            reader.close();
            // Using Prepared Statements (More secure)
            String selectTemplate = "SELECT colName FROM tName WHERE col = ?";
            PreparedStatement pstmt = conn.prepareStatement(selectTemplate);
            pstmt.setString(1, parameter); // Also setInt, setDate, ....
            ResultSet rset = pstmt.executeQuery(); // process rset, then close resources

            int numRowsAffected = stmt.executeUpdate(); // Inserting needs different exec

            conn.close(); // Close the connection
        } catch (SQLException e) { // Something was wrong with the SQL
            return;
        }
    }
}

```